

Search and Retrieval Algorithms

Ryan Layer

May 7, 2025

1 Introduction

Searching for a value in a list is a fundamental operation that many algorithms can solve, each with unique performance characteristics. These distinct features determine the suitability of each algorithm for specific problem types. Here, we explore the empirical time and memory tradeoffs of three popular options: binary search, merge search, and hash table search. Binary search sorts the database and then uses binary search to find each query in the database ($O(D \log D) + O(Q \log D)$). Merge search sorts both lists and then uses the merge operation to step through the lists in order to find matches ($O(D \log D) + O(Q \log Q) + O(Q + D)$). The hash table search places each database element in the database into a hash table and then looks for each query ($O(D) + O(Q)$, amortized). As expected, hash tables were the fastest and used the most memory, making them ideal for scenarios where database size is not a constraint. Notably, merge search outperformed binary search in handling large query sets, and should be used when the query set constitutes at least 50% of the database size.

2 Results

As expected, the hash table search was between 2X and 5X faster than the other methods but required over 4X more memory (Figure 1). Binary and merge search had identical memory footprints and similar runtime across the query set size range, with binary search slightly faster than merge search for small query sets and merge search slightly faster for larger query sets. Theoretically, we did not expect merge search to be faster than binary search. Memory usage was flat for all methods, indicating that the constant database size was the most significant factor.

Figure 1:




Figure 2: Enter Caption

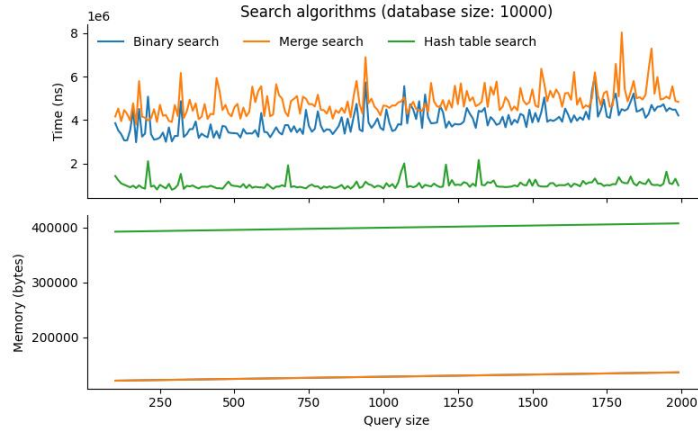


Figure 3: Chart with database size of 10000 and queries remaining at 100 only sampled once

3 Methods

3.1 Empirical comparison

We measured the time and memory usage of binary, merge, and hash table searches considering a database of 200,000 strings and query sets ranging from 100 to 200,000 strings. Strings for the database and query sets were drawn randomly from a word list without replacement. We created a new database and query set for each query set size and ran each search method, recording the run time and memory usage separately. We repeated this step five times and retained the meantime and memory metrics.

3.2 Theoretical comparison

To compare the theoretical performance of binary search and merge search, we used $\mathcal{T}(D, Q) = O(D \log D) + O(Q \log D)$ as the theoretical runtime function for binary search and $\mathcal{T}(D, Q) = O(D \log D) + O(Q \log Q) + O(Q + D)$ as the theoretical runtime for merge search, where D is the database size and Q is the query set size. We then set $D = 200,000$, let Q vary from 100 to 200,000, and plotted the runtime.

3.3 Reproducibility

To replicate these experiments, clone the repository and then run the empirical and simulation experiments as follows:

```
$ clone https://github.com/ryanlayerlab/search.git
$ cd search
$ python src/search.py \
  --query_range 100 200000 10000 \
  --database_size 200000 \
  --rounds 5 \
  --values_file data/words.txt.gz \
  --out_file doc/q100-10000_d200000_str.png
$ python sim.py \
  --query_range 100 200000 10000 \
  --database_size 200000 \
  --width 5 \
  --height 4 \
  --out_file ../doc/q100-10000_d200000_sim.png
```